

Lab 3 Concurrency

Deadlines & Grading

- **Prelab: (5pts)** Tues March 12 at 11:59 PM on Canvas only *one* team member needs to submit.
 - **Code and report submission:** (90pts) Friday, March 15 at **11:59 PM** on CMS, **groups of 2**
 - **Team Evaluation: (5pts)** Wed, March 20st **11:59 PM** on Canvas, individual
 - **Grading:** 9% of total grade
-

Prelab Questions (5pts)

To help you clarify your tasks and clear up any lingering questions about the material, this lab has a pre-lab. This lab is more complex than the previous two labs. In order to prepare, and discover some of the pitfalls, try to start the conversation with your new lab partner by explaining the following points to each-other. Only one of you needs to fill out the pre-lab.

Step 1: Please read the entire handout! This is the time to clarify anything with the course staff.

Step 2: Please answer the following questions in 1-3 sentences each. Use your own words. You may want to type out the answers in a separate document and then cut/paste them into the submission form.

Question 1: Were you able to make contact and meet with your lab partner? Hopefully this is a yes. If not what did you try to make this happen?

Question 2: Read through the support code. For the assembly code, it might make sense to print it out and use a high-lighter to trace the control flow of the assembly code on how/when the different portions get called. Describe what happens when a context switch occurs? What inputs/variables determine the behavior of `process_select`, and what the possible combinations your function needs to handle?

Question 3: It's never too early to think about testing! Describe how you will test the functionality of our multi-processing system processes that might be challenging as corner cases. What are they? How will you decide on the stack size for your processes?

Section I: Overview

Goal. The purpose of this lab is to give you experience with concurrency on the FRDM-KL46Z microcontroller. To successfully complete this lab, you will need to understand the following:

- I/O and Interrupts (Lab 2)
- Concurrency (time sharing) and its implementation

NOTE 1

We suggest that you first understand context switching on paper before you start writing C code. It is important that you have confidence in your implementation before you start, because it is challenging to find mistakes using standard testing methods.

Precautions.

- The micro-controller boards should be handled with care. Misuse such as incorrectly connecting the boards is likely to damage the device.
- These devices are *static sensitive*. This means that you can "zap" them with static electricity (a bigger problem in the winter months). Be very careful about handling boards that are not in their package. Your body should be at the same potential as the boards to avoid damaging them. For more information, check [wikipedia](https://en.wikipedia.org/wiki/Static_electricity).
- It is your responsibility to ensure that the boards are returned in the same condition as you received them. If you damage the boards, it is your responsibility to get a replacement.

Section II: Getting Started

Use the same setup instructions as in Lab 1 and Lab 2. Delete the automatically generated C-files and include the following provided files in your source.

3140.s – contains the assembly language definition of the timer interrupt and functions called when a process terminates.

3140_concurr.c – contains C definitions of functions for allocating stack space, initializing a process, and manipulating the process queue.

3140_concur.h – a header file with all function definitions listed (**including the ones that you must implement**).

process.h – header file for process type definitions. (**You need to modify this file**).

led.h – header file for helper functions for setting up and using the LEDs.

led.c/led_lowlevel.s – implementations of the helper functions.

test_1.c – a test program that uses concurrency using the on-board leds.

test_2.c – a test program that uses concurrency to pulse the serial LEDs on the board.

Add the provided source files to your project by directly extracting them from the zip file and moving them to the “source” folder for your project. Alternatively, you can also extract them elsewhere on your computer, right click on the “source” folder under your project, choose Import->General->Filesystem, choose the folder that the files were extracted to, and then select the files and click Finish.

Create a new file *process.c* which will implement the functions needed for the lab. You can create a new file by right clicking *source* in the Project window and clicking on *New -> File* Dialog. You should add the following includes to the top of your file:

```
#include "3140_concur.h"  
#include <stdlib.h>  
#include <MKL46Z4.h>
```

The *stdlib.h* header file allows you to use the `malloc` and `free` functions. The *3140_concur.h* header file gives the declarations for the functions you will need to implement. The *MKL46Z4.h* header file allows you to interact with peripherals on the device as in Lab 2.

NOTE 2

Look through all the provided files to make sure you understand the code provided. The provided files should not be modified. Instead, all the functionality you will need to implement should be written in the file *process.c*.

Section III: Assignment

Part 1: Implementing Concurrency

For this part of the lab assignment, you must implement three functions. These functions are directly called by users of your concurrency package and their functionality will be described below. These functions will manipulate the process queue, which is implemented as a linked list. Look at the appendix for more information.

```
int process_create (void (*f)(void), int n);
```

This function creates a process that starts at function `f`, with an initial stack size of `n`. It should return `-1` on an error, and `0` on success. The implementation of this function may require that you allocate memory for a `process_t` structure, and you can use `malloc()` for this purpose.

The state of a process can be initialized by calling the provided function `process_stack_init` (you can find the header of this function in *3140_concur.h*) which

allocates a new stack of size at $n \times 4$ bytes of free space in addition to an initialization context for the process. It returns the value of the stack pointer. This function returns NULL if the stack could not be allocated.

```
void process_start (void);
```

As discussed in class, a context switch occurs on a timeout that is triggered by a timer interrupt. For this lab, we will use the periodic interval timer described in Lab 2. As such, `process_start()` must:

1. Set up a period interrupt to be generated using the PIT interrupt. You will also need to call `NVIC_EnableIRQ` here. The interrupt handler itself is provided for you in *3140.s*, so you do NOT need to provide a C implementation of it.
2. Make sure all data structures you need are correctly initialized.
3. End with a call to `process_begin()`, which initiates the first process.

The context switches should occur at a rate of about 100Hz.

```
unsigned int * process_select(unsigned int * cursp);
```

This function is called by our provided (PIT0) interrupt handler code to select the next ready process. `cursp` will be the current stack pointer for the process that was just pre-empted and has the context at the top of the stack (NULL if there is no currently running process). If there are not processes left to run, this `process_select` should return NULL; if there is a ready process, this must return the value of the stack pointer for the next ready process. Your implementation must always maintain the global variable:

```
process_t *current_process_p
```

that points to currently running process. This variable should be set to NULL when the last process terminates, and should also be NULL until `process_start()` is called.

NOTE 3 You must use good programming practice and free memory that is no longer being used. This includes linked list nodes and process stacks for terminated processes. Linked list nodes can be freed using the function `free`. Process stacks can be freed using the function `process_stack_free` which is defined in *3140_concur.h* and implemented in *3140_concur.c*.

Task

Implement the three functions (`process_create`, `process_start`, and `process_select`) described above in *process.c*.

Part 2: Testing your Implementation

We have provided a test case (*test_0.c*) for your use as a basic check. Your lab should pass this test case, but it is not comprehensive enough to ensure that your code is functionally correct. You should design your own test cases in order to test the functionality.

Task

You must provide *two* additional tests to demonstrate that your implementation works correctly. Each test case should focus on some non-trivial aspect of the scheduler's behavior. Some examples are:

- What happens if there is only one call to `process_create()`?
- What if a process exits immediately?
- What if there are many short processes and one long process?

You must come up with at least one original test case (not one of the above). Your test cases should be documented (commented) to explain what they are testing along with a description of expected behavior. Name these files *mytest_1.c* and *mytest_2.c*.

Section IV: Submission

The lab requires the following files to be submitted:

- **process.h/process.c:** (50 pts)
You do not need to submit the support files we provided. Your implementation should be entirely contained in *process.h* and *process.c*. Naturally, your code should not depend on any changes in the support files.
- **test_1.c and test_2.c:** (20 pts)
These two files should have a main function and call *process.c* similar to the provide test case. The scoring will depend on the content or the report (testing section) as well.
- **report.pdf:** (10 pts + 10 style)
Provide a detailed description of your implementation of processes (max 5 pages in 11pt font and single line spacing – we will not read anything after page 5).

We recommend using illustrations to describe your key data structures (processes, etc.) and key properties of the implementation that you use to ensure correct operation. This write-up should also demonstrate that you understand the files that were provided (*3140.s*, *3140_concur.c*) in addition to your own implementation.

- Design
 - A high-level overview of your code. Did you make and design choices on the data-structure? Explain what the provided files do for you, and what they expect in terms of the different process stacks.
- Coding

- Include any specific details involved in implementing your design. Did you have to make any assumptions, i.e. number of processes, maximum stack use by each process?
- Code Review and Testing.
 - If you followed well-documented techniques (e.g., peer programming, naming conventions, etc.), this is the place to describe it.
 - Describe your testing strategy and how you ensured correct code functionality. For each test case, what did you mean to test in your test and how did you design the test? Where there any things that were difficult to test for?
- Work Distribution
 - How did you collaborate? In your own words, explain how you divided the work, how you communicated with each other, and whether/how everyone on the team had an opportunity to play an active role in all the major tasks. If you used any sort of tools to facilitate collaboration, describe them here. If there were any issues with the team, how did you try to resolve them? Note, you will also get a chance to give individual feedback and rate the collaboration, but we would like you to document what you intended to do here.
- Index
 - Please include any drawings or diagrams to support your work. This section will not count towards the 5-page limit. This layout helps us ensure the page limit and timely grading.

We encourage you to use meaningful variable and function names, comments, etc. to enhance code comprehensibility.

All files should be uploaded to CMS before the deadline. Multiple submissions are allowed, but only the latest submission will be graded.

Appendix

Data Structures

Process State

You will need a data structure to maintain the process state, and the linked-list based queue we provide expects the from detailed in `process.h`. The process type should contain both the data necessary to manage the process and pointer to the next process.

```
struct process_state {
    unsigned int * sp;
    /* the stack pointer for the process */
    ...
};
```

A process may end up relinquishing the processor to allow another process to execute at any time. We will achieve this by using timer interrupts. However, it may be advisable to disable these interrupts temporarily inside a process to create larger atomic operations and ensure that it does not get switched out of the processor. The functions that support disabling/enabling the interrupts (while keeping the timer running) respectively, are:

```
PIT->CHANNEL[0].TCTRL = PIT_TCTRL_TEN_MASK
PIT->CHANNEL[0].TCTRL = PIT_TCTRL_TEN_MASK | PIT_TCTRL_TIE_MASK;
```

Alternatively, `test_2.c` shows how to en/dis-able interrupts globally, and why it is important to create atomic blocks for correct program operation.

When a process terminates, the next process should be automatically selected for execution by calling your `process_select()` function. The support code will make the function call, but you need to handle this event in `process_select()`.

Linked Lists

Linked lists queues and stacks. They can easily grow and shrink, and they also make useful tasks for more complex operations, e.g. insertion, easy. A simple linked list node that holds an integer is given by, note the similarity to how we defined `process_t`:

```
struct mylist {
    int val;
    struct mylist *next;
};
```

You can use a single pointer to store the beginning of a linked list. Suppose this pointer is called `list_start`.

Initially, the list is empty. We need a special value to indicate pointers that point to nothing. In C, a common practice is to use `NULL` (integer 0) to indicate it. Hence, an empty list would be declared and initialized by

```
struct mylist * list_start = NULL;
```

If `elem` is a pointer to another `struct mylist` element, then we can insert it at the beginning of our list by the following operations:

```
elem->next = list_start;
list_start = elem;
```

The operation `elem->next` is shorthand for `(*elem).next`. We can traverse the list one item at a time as follows:

```
struct mylist *tmp;
for (tmp = list_start; tmp != NULL; tmp = tmp->next) {
    // tmp->val is the integer of interest
    // do something here
}
```

To insert an element at the end of the list is a bit more complicated. This is because there are two cases: (i) the list is currently empty; or (ii) the list has some items in it. If the list is empty, then the operation is easy. If the list is not empty, we need to traverse the list to find the last element in it, and then add the new one to the end of the list. The following code does this (assuming that `elem` is the new element to be added to the list):

```
struct mylist *tmp;
if (list_start == NULL) {
    list_start = elem;
    elem->next = NULL;
}
else {
    tmp = list_start;
    while (tmp->next != NULL) {
        // while there are more elements in the list
        tmp = tmp->next;
    }
    // now tmp is the last element in the list
    tmp->next = elem;
    elem->next = NULL;
}
```


We can remove the first element of the list in a straightforward way:

```
if (list_start == NULL) {
    elem = NULL;
}
else {
    elem = list_start;
    list_start = list_start->next;
}
```

In the above code, `elem` is the first element, and `elem->val` is the value. Normally we do not use `elem->next` at this point since we have removed it from the list. Sometimes programmers set `elem->next = NULL` just to be safe. A similar technique can be used to remove the last element of the list. While you can use this basic version of a linked list, a useful extension would be a slightly more complex list that keeps track of both the start and end, so that you can quickly insert tasks at the end without having to traverse the list firsts.

Bringing this back our problem. We implement insertion and removal from a process queue for you by providing `enqueue` and `dequeue` that operate on a type `process_queue_t` (which is essentially just the start pointer). Having a separate type for queues makes the code much easier to read since then the nodes and the object that holds them are different types. We also provide a helper function to check if the queue is empty. Even though we provide functionality for managing the queue, you still need to define the actual content of the `process_t` type.